

The ILOC Simulator

User Documentation

COMP 412, Fall 2015

Documentation for Lab 1

The **ILOC** instruction set is taken from the book, *Engineering A Compiler*, published by the Elsevier Morgan-Kaufmann [1]. The simulator itself was written over a period of many years by Tim Harvey, Todd Waterman, Keith Cooper, and, perhaps, others. The intent was to provide a useful tool for programming exercises in Rice's compiler courses. The author's intent was always that this software be distributable without cost for educational use.

1 Introduction and Roadmap

The **ILOC** simulator implements a subset of the **ILOC** operations described in Appendix A of *Engineering a Compiler, Second Edition* [1]. The simulator takes as input a file of **ILOC** operations. It simulates the execution of those operations. It reports on the results of that execution.

The simulator was designed as a target for the programming exercises in compiler courses at Rice, specifically the local register allocator and local instruction scheduler labs in the introductory course, COMP 412, and the optimizer lab in the scalar optimization course, COMP 512. The simulator has internal configuration parameters that change the number and type of operations allowed in a given cycle, that change the latencies of operations, and that change the set of allowed **ILOC** operations. These parameters allow one code base to generate the simulator for multiple programming assignments, with minimal source-level changes.

Roadmap This document describes the operation of the simulator. Section 2 describes the various command-line options that can control the simulator's behavior. Section 3 describes the **ILOC** virtual machine, the operations that it supports, and their formats and effects. Section 4 provides a brief discussion of the three mechanisms for initializing data memory in the **ILOC** virtual machine. Section 6 describes the execution tracing facility and the information that it provides. Finally, Section 7 describes the **ILOC** virtual machine configurations supported for different programming assignments, including machine parameters and operation latencies.

2 Command-Line Syntax

Command-line flags and arguments control the simulator's behavior. The command-line syntax for `sim` is as follows:

```
sim [options] [filename]
```

where the options are drawn from:

<code>-h</code>	Prints a list of all the command-line options,
<code>-d</code>	Specifies a <i>data initialization file</i> (see § 4).
<code>-m NUM</code>	Sets the number of <i>bytes</i> of available data memory to <code>NUM</code> .
<code>-r NUM</code>	Sets the number of available registers to <code>NUM</code> .
<code>-s NUM</code>	Sets the simulator's stall mode based on the value of <code>NUM</code> Mode 0 has no interlocks. Mode 1 interlocks on branches. Mode 2 interlocks on branches & memory. Mode 3 interlocks on branches, memory, & registers. The default stall mode is 3.
<code>-t</code>	Turns on execution tracing, which prints a record of each operation. Trace output is explained in § 6.
<code>-x</code>	Comp 412, Fall 2015, Lab 1 Code Check This option reports FAILURE if any register is defined more than once.
<code>-i NUM₀ NUM₁ ... NUM_n</code>	Treats <code>NUM₁</code> through <code>NUM_n</code> as integer data items and initializes data memory by writing them into consecutive words of memory, starting at the address <code>NUM₀</code> .
<code>-c NUM₀ NUM₁ ... NUM_n</code>	Operates in the same manner as <code>-i</code> , except that it treats <code>NUM₁</code> through <code>NUM_n</code> as bytes.

If `filename` is specified on the command line, the simulator expects that `filename` is a valid **ILOC** source file. It will read the file and execute the code, starting at the first operation in the code. If `filename` is not specified, the simulator reads from the standard input stream, `stdin` on Unix-like systems.

The `-m` and `-r` options must occur before any `-i` or `-c` option. Either `-i` and `-c` will cause the simulator to allocate its data structures so that it can initialize them. Once the data structures have been allocated, the `-m` and `-r` options have no effect.

The `-h`, `-d`, `-m`, `-r`, `-s`, and `-t` options should occur only once on the command line. Both the `-i` and `-c` options can occur multiple times, to initialize different address ranges in memory.

The operation and use of the `-d` option is explained in § 4.

3 The ILOC Virtual Machine

The simulator implements a simple virtual machine that supports a subset of **ILOC** [1, App. A]. Some properties of the virtual machine are fixed. For example, it features separate address spaces for code and data; a running program cannot read or write code memory. Some properties of the virtual machine are configurable from the command line. For example, the size of data memory, the number of registers, and the set of interlocks enforced on operations (see § 5) are all governed by command-line flags. Finally, some properties are determined when the simulator is, itself, compiled. For example, restrictions on the number of operations that execute per cycle (roughly speaking, the number of functional units that it emulates) and operation latencies are easily configurable. Section 7 describes the configurations used for specific programming assignments in Rice course.

The **ILOC** operations supported in the simulator fall into four basic categories: computational operations, data movement operations, control-flow operations, and output operations. Tables 1, 2, and 3 describe each group of operations. Note that **ILOC** is case sensitive. The 'I' used in immediate operations is the only uppercase letter.

An **ILOC** instruction is either a single operation, or a group of operations enclosed in square brackets and separated by semicolons, as in [*op*₁ ; *op*₂]. An instruction label in **ILOC**

Opcode	Format	Meaning
nop	nop	no change to registers or memory used as a placeholder or to cause a delay
add	add $r_0, r_1 \Rightarrow r_2$	$r_2 \leftarrow r_0 + r_1$
addl	addl $r_0, c_1 \Rightarrow r_2$	$r_2 \leftarrow r_0 + c_1$
sub	sub $r_0, r_1 \Rightarrow r_2$	$r_2 \leftarrow r_0 - r_1$
subl	subl $r_0, c_1 \Rightarrow r_2$	$r_2 \leftarrow r_0 - c_1$
mult	mult $r_0, r_1 \Rightarrow r_2$	$r_2 \leftarrow r_0 \times r_1$
multl	multl $r_0, c_1 \Rightarrow r_2$	$r_2 \leftarrow r_0 \times c_1$
div	div $r_0, r_1 \Rightarrow r_2$	$r_2 \leftarrow r_0 \div r_1$
divl	divl $r_0, c_1 \Rightarrow r_2$	$r_2 \leftarrow r_0 \div c_1$
lshift	lshift $r_0, r_1 \Rightarrow r_2$	$r_2 \leftarrow r_0 \ll r_1$
lshifl	lshifl $r_0, c_1 \Rightarrow r_2$	$r_2 \leftarrow r_0 \ll c_1$
rshift	rshift $r_0, r_1 \Rightarrow r_2$	$r_2 \leftarrow r_0 \gg r_1$
rshifl	rshifl $r_0, c_1 \Rightarrow r_2$	$r_2 \leftarrow r_0 \gg c_1$
and	and $r_0, r_1 \Rightarrow r_2$	$r_2 \leftarrow r_0 \&\& r_1$ (logical and)
andl	andl $r_0, c_1 \Rightarrow r_2$	$r_2 \leftarrow r_0 \&\& c_1$ (logical and)
or	or $r_0, r_1 \Rightarrow r_2$	$r_2 \leftarrow r_0 \parallel r_1$ (logical or)
orl	orl $r_0, c_1 \Rightarrow r_2$	$r_2 \leftarrow r_0 \parallel c_1$ (logical or)
not	not $r_0 \Rightarrow r_1$	$r_1 \leftarrow ! r_0$ (logical complement)

Table 1: The ILOC Computational Operations

Opcode		Format		Meaning
loadl	loadl	c_0	$\Rightarrow r_1$	$r_1 \leftarrow c_0$
load	load	r_0	$\Rightarrow r_1$	$r_1 \leftarrow \text{WORD}[r_0]$
loadAl	loadAl	r_0, c_1	$\Rightarrow r_2$	$r_1 \leftarrow \text{WORD}[r_0 + c_1]$
loadAO	loadAl	r_0, r_1	$\Rightarrow r_2$	$r_1 \leftarrow \text{WORD}[r_0 + r_1]$
cload	cload	r_0	$\Rightarrow r_1$	$r_1 \leftarrow \text{BYTE}[r_0]$
cloadAl	cloadAl	r_0, c_1	$\Rightarrow r_2$	$r_1 \leftarrow \text{BYTE}[r_0 + c_1]$
cloadAO	cloadAl	r_0, r_1	$\Rightarrow r_2$	$r_1 \leftarrow \text{BYTE}[r_0 + r_1]$
store	store	r_0	$\Rightarrow r_1$	$\text{WORD}[r_1] \leftarrow r_0$
storeAl	store	r_0	$\Rightarrow r_1, c_2$	$\text{WORD}[r_1 + c_2] \leftarrow r_0$
storeAO	store	r_0	$\Rightarrow r_1, r_2$	$\text{WORD}[r_1 + r_2] \leftarrow r_0$
cstore	cstore	r_0	$\Rightarrow r_1$	$\text{BYTE}[r_1] \leftarrow r_0$
cstoreAl	cstore	r_0	$\Rightarrow r_1, c_2$	$\text{BYTE}[r_1 + c_2] \leftarrow r_0$
cstoreAO	cstore	r_0	$\Rightarrow r_1, r_2$	$\text{BYTE}[r_1 + r_2] \leftarrow r_0$
i2i	i2i	r_0	$\Rightarrow r_1$	$r_1 \leftarrow r_0$, as an integer
c2c	c2c	r_0	$\Rightarrow r_1$	$r_1 \leftarrow r_0$, as a character
i2c	i2c	r_0	$\Rightarrow r_1$	$r_1 \leftarrow r_0$, as a character
c2i	c2i	r_0	$\Rightarrow r_1$	$r_1 \leftarrow r_0$, as a character

Table 2: The ILOC Data-Movement Operations

consists of an alphabetic character followed by zero or more alphanumeric characters. Any **ILOC** instruction may be labeled; the label precedes the instruction and is followed by a colon, as in L01: add r1, r2 \Rightarrow r3, or L02: [add r1, r2 \Rightarrow r3 ; i2i r0 \Rightarrow r4].

In the tables, r_i represents a register name; the subscripts make explicit the correspondence between operands in the “Format” column and the “Meaning” column. The notation c_i represents an integer constant, and L_i represents a label. “WORD[ex]” indicates the contents of the word of data memory at the location specified by ex . The address expression, ex , must be *word-aligned*—that is ($ex \bmod 4$) must be 0. “BYTE[ex]” indicates the contents of the byte of data memory at the location specified by ex , without an alignment constraint.

Opcodes are spelled as in the tables. They are case sensitive. The assignment arrow, drawn as \Rightarrow in the tables, is expressed in the code as an equals sign (=) followed by a greater than sign (>); the result is ‘=>’. The control-transfer arrow, drawn as \rightarrow in the tables, is expressed as a hyphen (-) followed by a greater than sign (>); the result is ‘->’.

Register names have an initial **r** followed immediately by a non-negative integer. The ‘r’ is case sensitive. Leading zeroes in the register name are not significant; thus r017 and r17 refer to the same register. Arguments that do not begin with an ‘r’ which appear as a c in the tables, are assumed to be positive integer constants in the range 0 to $2^{31}-1$.

Blanks and tabs are treated as whitespace. All **ILOC** opcodes must be followed by whitespace—any combination of blanks or tabs. Whitespace preceding and following other

Opcode	Format		Meaning
br	br	$\rightarrow L_0$	control transfers to L_0
cbr	cbr	$r_1 \rightarrow L_1, L_2$	if r_0 is true, control transfers to L_1 otherwise, control transfers to L_2
cmp_LT	cmp_LT	$r_0, r_1 \Rightarrow r_2$	$r_2 \leftarrow \text{true}$ if $r_0 < r_1$ otherwise, $r_2 \leftarrow \text{false}$
cmp_LE	cmp_LE	$r_0, r_1 \Rightarrow r_2$	$r_2 \leftarrow \text{true}$ if $r_0 \leq r_1$ otherwise, $r_2 \leftarrow \text{false}$
cmp_GT	cmp_GT	$r_0, r_1 \Rightarrow r_2$	$r_2 \leftarrow \text{true}$ if $r_0 > r_1$ otherwise, $r_2 \leftarrow \text{false}$
cmp_GE	cmp_GE	$r_0, r_1 \Rightarrow r_2$	$r_2 \leftarrow \text{true}$ if $r_0 \geq r_1$ otherwise, $r_2 \leftarrow \text{false}$
cmp_EQ	cmp_EQ	$r_0, r_1 \Rightarrow r_2$	$r_2 \leftarrow \text{true}$ if $r_0 = r_1$ otherwise, $r_2 \leftarrow \text{false}$
cmp_NE	cmp_NE	$r_0, r_1 \Rightarrow r_2$	$r_2 \leftarrow \text{true}$ if $r_0 \neq r_1$ otherwise, $r_2 \leftarrow \text{false}$
output	output	c_0	writes $\text{WORD}[c_0]$ to <code>stdout</code> c_0 must be a word-aligned integer constant
coutput	coutput	c_0	writes $\text{BYTE}[c_0]$ to <code>stdout</code> as a character

Table 3: The ILOC Control-Flow and Output Operations

symbols is optional. Whitespace may not appear within operation names, register names, or the assignment symbol. A double slash (“//”) indicates that the rest of the line is a comment. Empty lines may appear in the input; the simulator will ignore them.

In addition to the **ILOC** operations, the simulator supports two pseudo-operations to initialize memory with known values; they are “pseudo” operations in the sense that their effects occur before the **ILOC** code starts execution.

Pseudo-Op	Format	Meaning
dcs	dcs c_0 c_1 c_2 $\dots c_n$	c_0 is an integer address. c_1 through c_n are characters separated by blanks. c_1 through c_n are placed in consecutive bytes of memory, starting at address c_0 .
dis	dis c_0 c_1 c_2 $\dots c_3$	c_0 is a word-aligned integer address. c_1 through c_n are integers. c_1 through c_n are placed in consecutive words of memory, starting at address c_0 .

Here, **dcs** is an acronym for *define character storage* and **dis** is an acronym for *define integer storage*. These pseudo-operations provide a convenient way to initialize large quantities of memory, such as an array.

The pseudo-operations must appear before the first **ILOC** operation. There is no arbitrary restriction on the number of pseudo-operations or their relative order. The obvious way to use these pseudo-operations is to create a “data” file and use the **-d** command-line flag to prepend it to an **ILOC** file. In this way, a single program can be easily run against multiple different data files.

4 Data-Initialization Methods

There are three mechanisms to initialize the simulator’s data memory before any code executed. The **-i** and **-c** command-line options allow the user to specify an address and a list of data values. The initialization occurs before any operation executes. Note that the command line may contain more than one **-i** or **-d** sequence. Negative numbers cannot occur in either of these sequences.

The other mechanism to initialize the simulator’s data memory is the **-d** option. This option lets the user specify a *data-initialization file* that contains one or more **dis** or **dcs** pseudo-operations, as described in §3. Only one data-initialization file can be specified.

The **-d** option has the following behavior. It creates a temporary file in the designated temporary file area (specified by the macro definition **TEMPDIR** in file **sim.h**). It copies the contents of the data initialization file to the temporary file. It then copies the contents of the code file, either **filename** from the command line or **stdin** if no **filename** is given. It closes the temporary file and reopens it as the scanner’s input file. After execution, it deletes the temporary file.

5 Interlocks

The **ILOC** simulator supports four distinct levels of enforced synchronization. These levels are:

- 3 interlocks on branches, registers, and memory
- 2 interlocks on branches and memory
- 1 interlocks on branches
- 0 no interlocks

The branch interlock prevents the simulator from executing operations during the latency of a branch. Once the branch operation is issued, the simulator will stall all pending operations until the branch completes.

The memory interlock stalls operations when a load or output operation tries to access a memory location that will be modified by a pending operation—one that has been issued and has not yet completed. Similarly, the register interlock stalls operations that attempt to read a value from a register that will be modified by a pending operation.

The interlock level is set by the **-s** command-line flag. The default interlock levels are documented in Section 7. The trace facility lists the interlock level in use for each run at the head of the trace.

6 Understanding the Trace Output

To help the user understand the details of a specific execution, the **ILOC** simulator includes a trace facility, invoked with the `-t` command-line flag. The trace shows each executed instruction, its arguments and its results.

Figure 1 shows an example trace generated by the **ILOC** code for the small program shown to the right. Below the source code is the **ILOC** code that the compiler generated. While the program is simple, it highlights several of the important features of the trace facility.

In the translation, both `a` and `i` are kept in registers. Thus, the **ILOC** code contains no `load` operations. The only `store` operation is generated by the `print` statement, which must store `a`'s value to memory so that it can generate an **ILOC** output operation. Notice that the code assumes a single functional-unit configuration of the simulator; it contains no explicitly specified instruction-level parallelism.

Turning our attention to the trace in Figure 1, it begins by listing the version number of the simulator and the interlock settings in use for the run. The final line in the trace is the execution summary generated by every simulator run, whether traced or not. Between the interlock settings and the execution summary, the trace lists the operations executed at each cycle of the execution, one cycle per line. Each line begins with the cycle number.

To the right of the cycle number, the operation(s) executed in that cycle are listed. For any register operand, the value of the register appears in parentheses after the register's name. Registers that are used show their values before the operation executes. Registers that are defined show their values after the operation takes effect.

For a long-latency operation, the result is shown in the trace for the cycle in which the instruction issues. When a long-latency operation completes, that fact is noted at the end of the trace for the cycle in which it completes. The trace for that cycle will show an asterisk (`*`) followed by the cycle number in which the long-latency operation was first issued.

Look at the `store` issued in cycle 26. Because the `output` uses the same memory location (and the simulator has memory interlocks enabled), the `output` operation stalls until the `store` completes in cycle 30. The trace for cycle 30 ends with the notation `*26` to indicate that an operation issued in cycle 26 completed at the end of cycle 30. The `output` is then issued in the next cycle.

Now, look at the `cbr` operations in cycles 9, 14, 19, and 24. The asterisk in the trace for `cbr` indicates which branch was taken. In the first three `cbr` operations, the branch transferred control to `L0` at the top of the loop. In the final `cbr`, it transferred control to `L1`, the label on the first statement after the loop.

When the simulator is generating a trace, it changes the format of the output generated by the `output` operation. In a run without tracing, the number is simply writted to the

```
procedure main {  
    int a, i;  
    a = 1;  
    for i = 1 to 4 by 1 {  
        a = a + 1;  
    }  
    print a;  
}
```

Source Program

```
loadI 1 => r0  
loadI 1 => r1  
loadI 4 => r2  
cmp_LE r1, r2 => r3  
cbr r3 -> L0, L1  
L0: addI r0, 1 => r4  
i2i r4 => r0  
addI r1, 1 => r1  
cmp_LE r1, r2 => r5  
cbr r5 -> L0, L1  
L1: loadI 0 => r6  
store r0 => r6  
output 0
```

ILOC Program

ILOC Simulator, Version 512-2-0
Interlock settings: memory registers branches

```
0: [loadI 1 => r0 (1)]
1: [loadI 1 => r1 (1)]
2: [loadI 4 => r2 (4)]
3: [cmp_LE r1 (1), r2 (4) => r3 (1)]
4: [cbr r3 (1) -> L0*, L1]
5: [addI r0 (1), 1 => r4 (2)]
6: [i2i r4 (2) => r0 (2)]
7: [addI r1 (1), 1 => r1 (2)]
8: [cmp_LE r1 (2), r2 (4) => r5 (1)]
9: [cbr r5 (1) -> L0*, L1]
10: [addI r0 (2), 1 => r4 (3)]
11: [i2i r4 (3) => r0 (3)]
12: [addI r1 (2), 1 => r1 (3)]
13: [cmp_LE r1 (3), r2 (4) => r5 (1)]
14: [cbr r5 (1) -> L0*, L1]
15: [addI r0 (3), 1 => r4 (4)]
16: [i2i r4 (4) => r0 (4)]
17: [addI r1 (3), 1 => r1 (4)]
18: [cmp_LE r1 (4), r2 (4) => r5 (1)]
19: [cbr r5 (1) -> L0*, L1]
20: [addI r0 (4), 1 => r4 (5)]
21: [i2i r4 (5) => r0 (5)]
22: [addI r1 (4), 1 => r1 (5)]
23: [cmp_LE r1 (5), r2 (4) => r5 (0)]
24: [cbr r5 (0) -> L0, L1*]
25: [loadI 0 => r6 (0)]
26: [store r0 (5) => r6 (addr: 0)]
27: [ stall ]
28: [ stall ]
29: [ stall ]
30: [ stall ] *26
31: [output 0 (5)]
output generates => 5
```

Executed 28 instructions and 28 operations in 32 cycles.

Figure 1: Execution Trace for Simple Example

standard output stream, one number per line. In a run with tracing, the simulator adds the text `output generates =>` to help the user find the printed result.

7 Details of Specific Simulator Configurations

This section provide simulator configuration details for specific assignments.

7.1 COMP 412, Fall 2015, Lab1

The simulator for Lab 1 supports a subset of the **ILOC** operations that consists of `load`, `loadI`, `store`, `add`, `sub`, `mult`, `lshift`, `rshift`, `output`, and `nop`. These operations are describes in Tables 1 and 2. All other operations are disabled.

- All operations have a single-cycle latency, except `load` and `store`, which have a three-cycle latency. That is, the value defined by a load or store operation is available at the start of the third cycle after the operation issues.¹
- The simulator is configured with a single functional unit. It issues one new operation per cycle.
- The simulator is configured with a default of 1,000,000 registers and 4,000,000 bytes of memory.
- The simulator is configured with a default interlock level of 3. It interlocks on memory and registers. Your lab should never encounter a branch operation.

The default values can all be changed at the command line, as described in § 2.

For Lab 1, the first 32 KB of memory is reserved for data used by the input block—that is, locations 0 through 32,767 are not available to your allocator for storage of spilled values. Memory starting at location 32,768 is available to store spilled values.

References

- [1] Keith Cooper and Linda Torczon. *Engineering A Compiler*. Elsevier Morgan Kaufmann Publishers Inc., San Francisco, CA, USA, 2nd edition, 2011.

¹You can see these values “commit” in the trace of a specific execution. For example, in Figure 1, the store issued in cycle 26 commits at the end of cycle 30. That particular trace was generated with a simulator configured for a load/store latency of five.